

Cow Face ID – End-to-End Model (YOLOv5 + ViT + ArcFace)

This repo-style guide gives you **ready-to-run code** to replicate the dairy-cow facial identification system described in your paper: **YOLOv5 for face detection + ViT-small with ArcFace** for identification. Copy the code blocks into files with the indicated filenames, follow the steps, and you'll have training + real-time inference.

0) Project Structure

```
cow-face-id/
├─ env.yml                # (optional) conda environment
├─ requirements.txt       # pip deps
├─ README.md             # (this file)
├─ data/
│   ├─ detection/        # YOLO dataset (images/, labels/)
│   │   ├─ images/{train,val,test}/
│   │   └─ labels/{train,val,test}/
│   └─ recognition/      # cropped faces (class folders) for ViT
│       ├─ train/<cow_id>/*.jpg
│       ├─ val/<cow_id>/*.jpg
│       └─ test/<cow_id>/*.jpg
├─ scripts/
│   ├─ build_crops_from_yolo.py
│   ├─ train_recognition_vit_arcface.py
│   ├─ infer_realtime.py
│   └─ utils_arcface.py
└─ yolo/                 # YOLOv5 repo will live here (git clone)
```

Tip: Put your raw barn videos/jpegs anywhere (e.g., `raw/`). Create YOLO labels for faces. Then run the crop builder to generate the recognition dataset.

1) Environment & Dependencies

Option A: Conda

```
# env.yml
name: cow-face-id
channels:
  - pytorch
  - nvidia
  - conda-forge
dependencies:
```

```
- python=3.10
- pytorch
- torchvision
- torchaudio
- pytorch-cuda=12.1 # if you have NVIDIA GPU
- pip
- pip:
  - -r requirements.txt
```

Pip requirements

```
# requirements.txt
opencv-python
numpy
pandas
scikit-learn
matplotlib
tqdm
albumentations
einops
torchmetrics
```

YOLOv5 is used via its own repo scripts for training/inference.

Install & clone YOLOv5:

```
conda env create -f env.yml # or: python -m venv .venv && source .venv/bin/
activate
conda activate cow-face-id

# clone YOLOv5 into ./yolo
cd yolo
git clone https://github.com/ultralytics/yolov5 .
pip install -r requirements.txt # of YOLOv5 repo
cd ..

# install this project's reqs
pip install -r requirements.txt
```

2) Prepare Data

2.1 YOLO Detection Dataset

- **Goal:** Detect cow faces in general frames.
- **Format:** Standard YOLO: one `.txt` per image with lines: `class x y w h` (normalized).
- **Split:** `train/`, `val/`, `test/` under both `images/` and `labels/`.

Example YAML for YOLOv5:

```
# data/detection/cow_face.yaml
path: ../data/detection
train: images/train
val: images/val
names:
  0: cow_face
```

2.2 Build Recognition Dataset (COW77-like)

Run the script below to **crop faces** from detection labels and build folders per cow id (you provide a CSV mapping image→cow_id).

Create file: `scripts/build_crops_from_yolo.py`

```
import os
import cv2
import csv
from pathlib import Path

IMAGES_ROOT = Path('data/detection/images') # source full frames
LABELS_ROOT = Path('data/detection/labels')
OUT_ROOT = Path('data/recognition')
SPLITS = ['train', 'val', 'test']

# CSV mapping for identity labels per frame: image_name,cow_id,split
# Example row: DSC_000123.jpg, COW_017, train
ID_CSV = Path('data/identity_map.csv')

# crops will be resized to 224x224 for ViT
TARGET = 224

os.makedirs(OUT_ROOT, exist_ok=True)

# helper to convert normalized xywh -> pixel xyxy
def xywhn_to_xyxy(label_path, img_w, img_h):
    boxes = []
    with open(label_path, 'r') as f:
        for line in f:
            parts = line.strip().split()
            if len(parts) != 5:
                continue
            cls, x, y, w, h = parts
            x, y, w, h = map(float, (x, y, w, h))
            x1 = int((x - w/2) * img_w)
            y1 = int((y - h/2) * img_h)
            x2 = int((x + w/2) * img_w)
            y2 = int((y + h/2) * img_h)
```

```

        boxes.append((x1, y1, x2, y2))
    return boxes

# read identity map
meta = {}
with open(ID_CSV, 'r', newline='') as f:
    reader = csv.DictReader(f)
    for r in reader:
        meta[r['image_name']] = (r['cow_id'], r['split'])

for split in SPLITS:
    img_dir = IMAGES_ROOT / split
    for img_name in os.listdir(img_dir):
        if not img_name.lower().endswith(('.jpg', '.jpeg', '.png')):
            continue
        if img_name not in meta:
            # skip frames not labeled with identity
            continue
        cow_id, split2 = meta[img_name]
        # allow override split from CSV
        split_out = split2 if split2 in SPLITS else split

        img_path = img_dir / img_name
        lbl_path = LABELS_ROOT / split / (Path(img_name).stem + '.txt')
        if not lbl_path.exists():
            continue

        img = cv2.imread(str(img_path))
        if img is None:
            continue
        h, w = img.shape[:2]
        boxes = xywhn_to_xyxy(lbl_path, w, h)
        if not boxes:
            continue

        # Pick the largest face box (closest cow) or loop over all
        for (x1,y1,x2,y2) in boxes:
            x1 = max(0, x1); y1 = max(0, y1)
            x2 = min(w-1, x2); y2 = min(h-1, y2)
            if x2 <= x1 or y2 <= y1:
                continue
            crop = img[y1:y2, x1:x2]
            crop = cv2.resize(crop, (TARGET, TARGET),
interpolation=cv2.INTER_AREA)
            out_dir = OUT_ROOT / split_out / cow_id
            out_dir.mkdir(parents=True, exist_ok=True)
            out_name = out_dir / f"{Path(img_name).stem}_{x1}_{y1}.jpg"
            cv2.imwrite(str(out_name), crop)

```

Run:

```
python scripts/build_crops_from_yolo.py
```

3) Train YOLOv5 Face Detector

From the `yolo/` folder you cloned:

```
cd yolo
python train.py \
  --img 640 \
  --batch 16 \
  --epochs 150 \
  --data ../data/detection/cow_face.yaml \
  --weights yolov5s.pt \
  --project runs/cow_face \
  --name y5s_detect
cd ..
```

This produces weights like `yolo/runs/cow_face/y5s_detect/weights/best.pt`.

You can later run detection to create more crops if needed:

```
cd yolo
python detect.py --weights runs/cow_face/y5s_detect/weights/best.pt \
  --img 640 --source ../data/detection/images/val --save-txt --save-conf
cd ..
```

4) Train ViT-small + ArcFace for Recognition

Create file: `scripts/utils_arcface.py`

```
import math
import torch
import torch.nn as nn

class L2Norm(nn.Module):
    def __init__(self, dim=1, eps=1e-12):
        super().__init__()
        self.dim = dim
        self.eps = eps
    def forward(self, x):
        return torch.nn.functional.normalize(x, p=2, dim=self.dim,
        eps=self.eps)
```

```

class ArcMarginProduct(nn.Module):
    """Implements ArcFace linear head.
    Args:
        in_features: embedding dim
        out_features: num classes
        s: scale factor
        m: additive angular margin
    """
    def __init__(self, in_features, out_features, s=10.0, m=0.25,
easy_margin=False):
        super().__init__()
        self.in_features = in_features
        self.out_features = out_features
        self.s = s
        self.m = m
        self.weight = nn.Parameter(torch.FloatTensor(out_features,
in_features))
        nn.init.xavier_uniform_(self.weight)
        self.easy_margin = easy_margin
        self.cos_m = math.cos(m)
        self.sin_m = math.sin(m)
        self.th = math.cos(math.pi - m)
        self.mm = math.sin(math.pi - m) * m

    def forward(self, embeddings, labels):
        # normalize features and weights
        embeddings = torch.nn.functional.normalize(embeddings, p=2, dim=1)
        W = torch.nn.functional.normalize(self.weight, p=2, dim=1)
        cosine = torch.matmul(embeddings, W.t())
        sine = torch.sqrt(1.0 - torch.clamp(cosine ** 2, 0, 1))
        phi = cosine * self.cos_m - sine * self.sin_m
        if self.easy_margin:
            phi = torch.where(cosine > 0, phi, cosine)
        else:
            phi = torch.where(cosine > self.th, phi, cosine - self.mm)
        one_hot = torch.zeros_like(cosine)
        one_hot.scatter_(1, labels.view(-1, 1), 1.0)
        logits = (one_hot * phi) + ((1.0 - one_hot) * cosine)
        logits *= self.s
        return logits

```

Create file: `scripts/train_recognition_vit_arcface.py`

```

import os
from pathlib import Path
import random
import numpy as np
from tqdm import tqdm

import torch

```

```

import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import transforms, datasets

from utils_arcface import ArcMarginProduct

# Config
DATA_ROOT = Path('data/recognition')
IMG_SIZE = 224
BATCH_SIZE = 64
EPOCHS = 150
LR = 0.05
MOMENTUM = 0.9
WD = 1e-4
SCALE = 10.0
MARGIN = 0.25
SEED = 42

random.seed(SEED); np.random.seed(SEED); torch.manual_seed(SEED)

def build_transforms():
    train_tf = transforms.Compose([
        transforms.Resize((IMG_SIZE, IMG_SIZE)),
        transforms.RandomApply([transforms.ColorJitter(0.2,0.2,0.2,0.1)],
p=0.5),
        transforms.RandomHorizontalFlip(p=0.2),
        transforms.RandomRotation(10),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224,
0.225])
    ])
    test_tf = transforms.Compose([
        transforms.Resize((IMG_SIZE, IMG_SIZE)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224,
0.225])
    ])
    return train_tf, test_tf

def create_loaders():
    train_tf, test_tf = build_transforms()
    train_ds = datasets.ImageFolder(DATA_ROOT / 'train', transform=train_tf)
    val_ds = datasets.ImageFolder(DATA_ROOT / 'val', transform=test_tf)
    test_ds = datasets.ImageFolder(DATA_ROOT / 'test', transform=test_tf)

    train_loader = DataLoader(train_ds, batch_size=BATCH_SIZE, shuffle=True,
num_workers=4, pin_memory=True)
    val_loader = DataLoader(val_ds, batch_size=BATCH_SIZE, shuffle=False,
num_workers=4, pin_memory=True)

```

```

    test_loader = DataLoader(test_ds, batch_size=BATCH_SIZE, shuffle=False,
num_workers=4, pin_memory=True)
    return train_loader, val_loader, test_loader, train_ds.classes

class ViTSmall(nn.Module):
    def __init__(self, embed_dim=384, num_classes=1000):
        super().__init__()
        # Use torchvision's ViT-S/16
        from torchvision.models.vision_transformer import vit_b_16,
ViT_B_16_Weights, vit_s_16, ViT_S_16_Weights
        self.backbone = vit_s_16(weights=ViT_S_16_Weights.IMAGENET1K_V1)
        # replace head with identity to output embeddings
        self.embedding_dim = self.backbone.heads.head.in_features
        self.backbone.heads.head = nn.Identity()

    def forward(self, x):
        return self.backbone(x)

def evaluate(model, head, loader, device):
    model.eval(); head.eval()
    correct = 0; total = 0
    with torch.no_grad():
        for x, y in loader:

```